

Experiments Using Variational

Lab 11 : Variational Autoencoder (VAE)

Aim: To implement a Variational Autoencoder (VAE) and study its generative ability to reconstruct

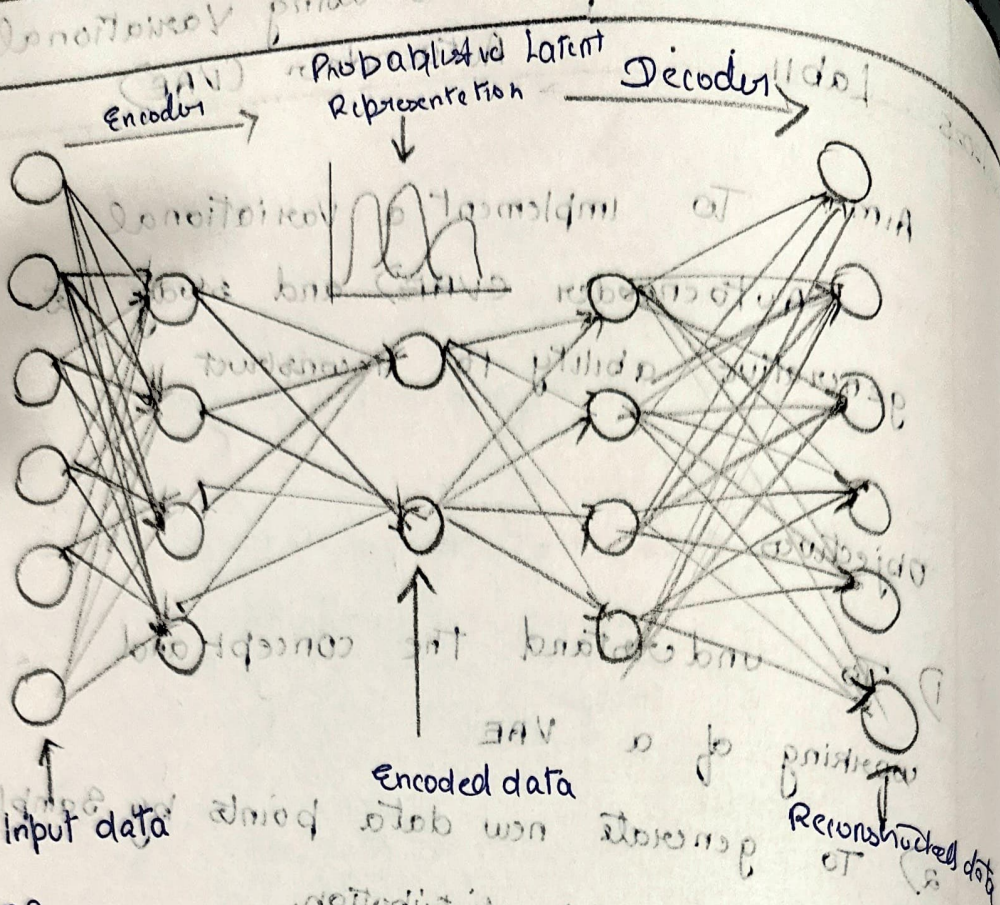
Objectives

- 1) To understand the concept and working of a VAE
- 2) To generate new data points by sampling from a latent distribution.
- 3) To visualize reconstructed & newly generated images
- 4) To compare the performance of VAE

Pseudo code

Start

- 1) Import libraries
- 2) Load MNIST dataset
- 3) Define encoder :
 - Input layer
 - Dense (256, activation = 'relu')
 - Dense (128, activation = 'relu')
 - output : mean (μ) and log variance (σ^2)
- 4) Sample latent
$$z = \mu + \sigma * \epsilon, \text{ where } \epsilon \sim N(0, 1)$$
- 5) Define Decoder :
 - Input



OP

Epoch [1/5], Loss: 164.0216

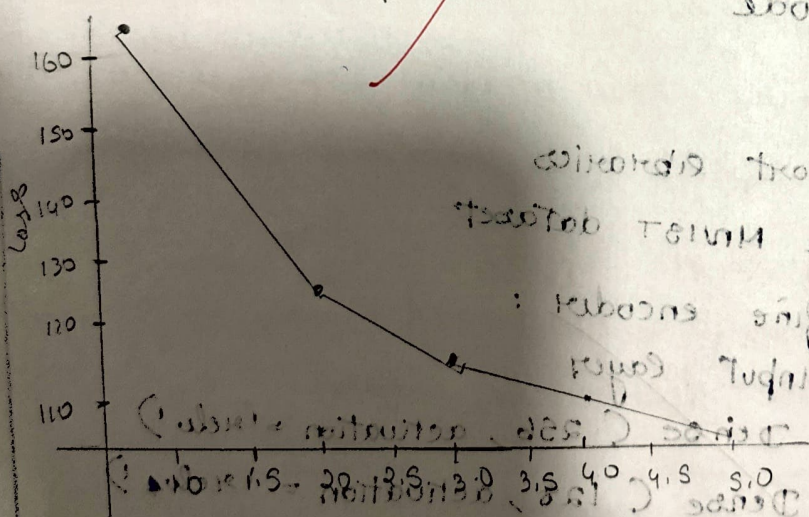
Epoch [2/5], Loss: 121.5716

Epoch [3/5], Loss: 114.6072

Epoch [4/5], Loss: 111.6099

Epoch [5/5], Loss: 109.8843

Epochs Loss



Epochs

6) Define Total loss :

$$\text{Total loss} = \text{Reconstruction loss} + \text{KL Divergence loss}$$

7) Compute and train model using MNIST dataset

8) Visualize

End

Observation - Fast

- 1) The autoencoder effectively reduced 784 dimensional input images into a 32-dimensional latent representation
- 2) The training loss decreased steadily with each epoch
- 3) Reconstructed images were visually similar to originals
- 4) The encoder captured essential patterns like stroke and shapes of handwritten digits
- 5) The model achieved efficient image compression.
- 6) The latent space representation could be used for feature extraction / dimensionality reduction in other ML tasks.

The VAE demonstrated better generative capability than traditional autoencoder.

Result

Successfully implemented VAE.

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
```

```
batch_size = 128
learning_rate = 1e-3
num_epochs = 5
latent_dim = 20
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
transform = transforms.ToTensor()
train_dataset = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
test_dataset = datasets.MNIST(root='./data', train=False, transform=transform, download=True)
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
```

```
class VAE(nn.Module):
    def __init__(self):
        super(VAE, self).__init__()
        # Encoder
        self.fc1 = nn.Linear(28*28, 400)
        self.fc_mu = nn.Linear(400, latent_dim)
        self.fc_logvar = nn.Linear(400, latent_dim)
        # Decoder
        self.fc3 = nn.Linear(latent_dim, 400)
        self.fc4 = nn.Linear(400, 28*28)

    def encode(self, x):
        h1 = torch.relu(self.fc1(x))
        return self.fc_mu(h1), self.fc_logvar(h1)

    def reparameterize(self, mu, logvar):
```



```
return self.decode(z), mu, logvar
```

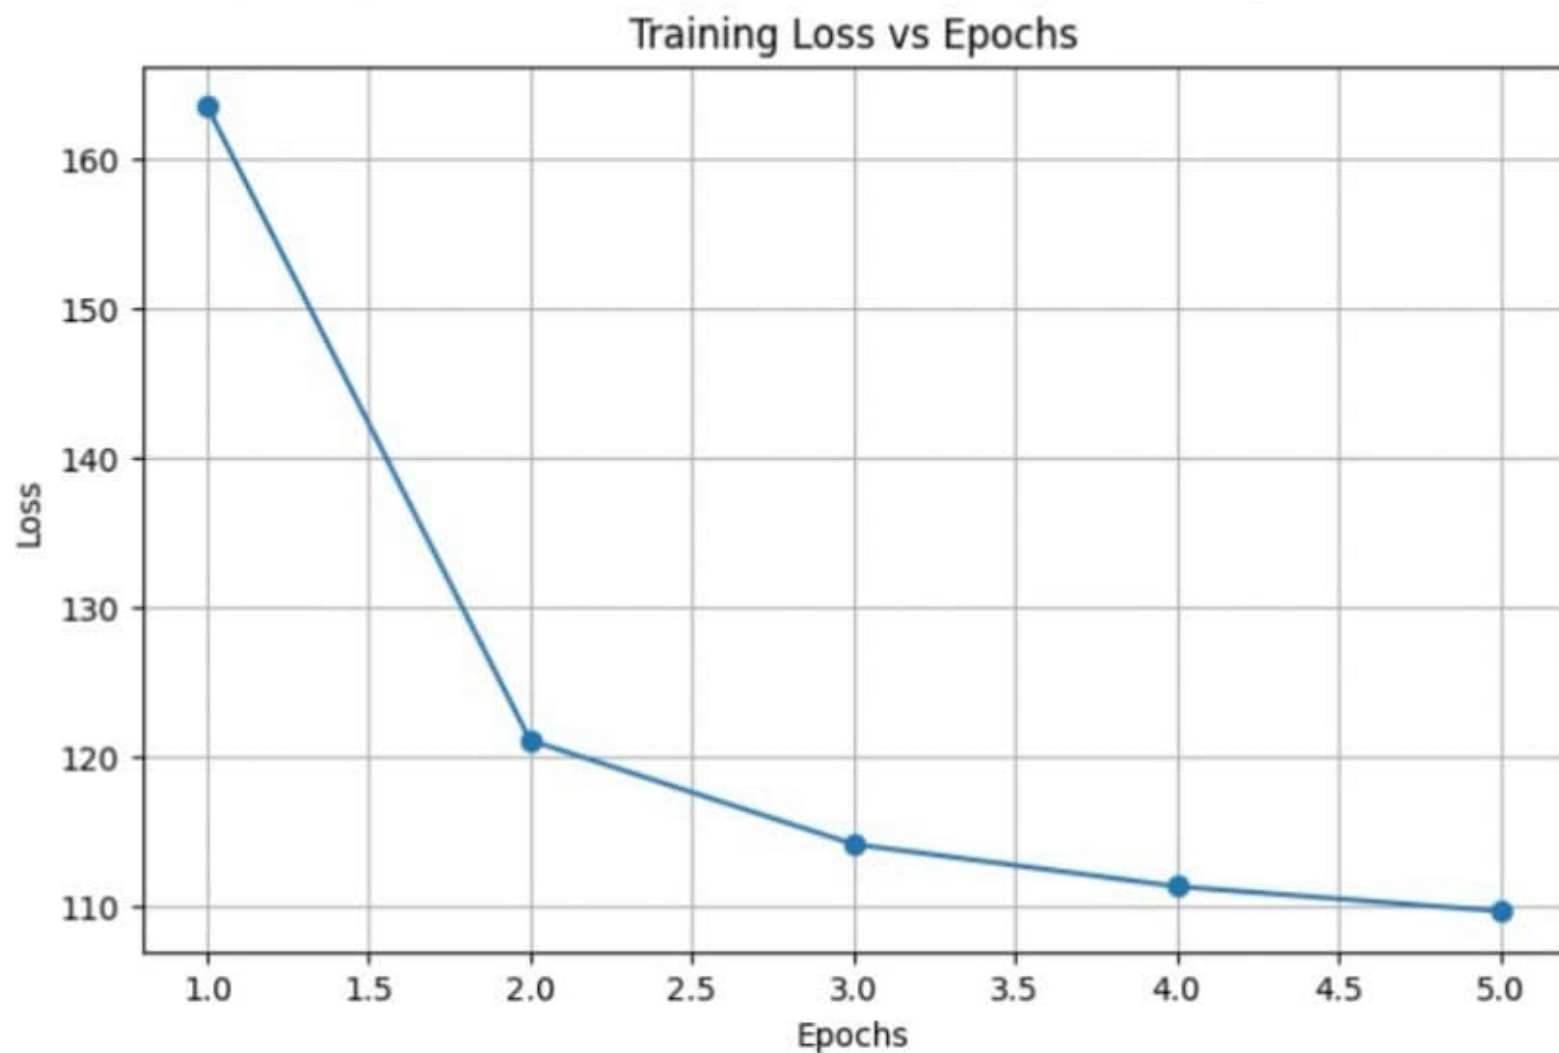
```
[15] def loss_function(recon_x, x, mu, logvar):  
    BCE = nn.functional.binary_cross_entropy(recon_x, x.view(-1, 28*28), reduction='sum')  
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())  
    return BCE + KLD
```

```
[16] model = VAE().to(device)  
optimizer = optim.Adam(model.parameters(), lr=learning_rate)
```

```
[17] train_losses = []  
  
for epoch in range(num_epochs):  
    model.train()  
    train_loss = 0  
    for data, _ in train_loader:  
        data = data.to(device)  
        optimizer.zero_grad()  
        recon_batch, mu, logvar = model(data)  
        loss = loss_function(recon_batch, data, mu, logvar)  
        loss.backward()  
        train_loss += loss.item()  
        optimizer.step()  
  
    avg_loss = train_loss / len(train_loader.dataset)  
    train_losses.append(avg_loss)  
    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}')
```

```
Epoch [1/5], Loss: 163.4689  
Epoch [2/5], Loss: 121.1113  
Epoch [3/5], Loss: 114.1747  
Epoch [4/5], Loss: 111.3527  
Epoch [5/5], Loss: 109.7027
```

```
plt.figure(figsize=(8,5))
plt.plot(range(1, num_epochs+1), train_losses, marker='o')
plt.title("Training Loss vs Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.grid(True)
plt.show()
```



```
plt.figure(figsize=(20,4))
for i in range(n):
    # Original
    ax = plt.subplot(2, n, i+1)
    plt.imshow(data[i].cpu().squeeze(), cmap='gray')
    plt.axis('off')
    # Reconstructed
    ax = plt.subplot(2, n, i+1+n)
    plt.imshow(recon[i].cpu().view(28,28), cmap='gray')
    plt.axis('off')
plt.suptitle("Top: Original | Bottom: Reconstructed", fontsize=14)
plt.show()
```

Top: Original | Bottom: Reconstructed

